


Graph Reachability

Panagiotis Parchas

HKUST



Where are the graphs?


- Some famous graph data
 - Social Networks
 - Web sites/internet
 - XML Documents
 - Biological/chemical networks

Why graph analysis?

- Knowledge discovery in social networks
 - Targeted advertising.
 - Pattern extraction for behavioral analysis.
 - Useful statistics and conclusions for a wide range of scientists
- Mining in Web Data
 - Interesting navigation patterns
 - XML documents
- Biological Data Analysis
 - Drug discovery
 - DNA Analysis
- Geographic Information Systems
-And much more.....

Problem formulation

Given a *directed graph* G and two nodes u and v , is there a path connecting u to v (denoted $u \rightsquigarrow v$)?



$3 \rightsquigarrow 11$? YES

$3 \rightsquigarrow 12$? NO

Directed Graph → DAG (directed acyclic graph) by
coalescing the strongly connected components

Motivation

- Classical problem in graph theory.
- Studying the influence flow in social networks
 - Even undirected graphs (facebook) are converted to directed w.r.t a certain attribute distribution
- Security: finding possible connections between suspects
- Biological data: is that protein involved- directly or indirectly- in the expression of a gene?
- Primitive for many graph related problems (pattern matching)

Methods Overview

	Method	Query time	Construction	Index size
Naïve	DFS/BFS	$O(n+m)$	$O(n+m)$	$O(n+m)$
	Transitive Closure	$O(1)$	$O(nm)=O(n^3)$	$O(n^2)$
Tree Cover	Optimal Tree Cover (Agrawal et al., SIGMOD'89)	$O(n)$	$O(nm)=O(n^3)$	$O(n^2)$
	GRIPP (Trißl et al., SIGMOD'07)	$O(m-n)$	$O(n+m)$	$O(n+m)$
	Dual-Labeling (Wang et al., ICDE'06)	$O(1)$	$O(n+m+t^3)$	$O(n+t^2)$
Chain Cover	Optimal Chain Cover (Jagadish, TODS'90)	$O(k)$	$O(nm)$	$O(nk)$
	Path-Tree (Jin, et al., SIGMOD'08)	$\log^2 k'$	$O(mk')/O(mn)$	$O(nk')$
HOP Cover	2-HOP (SODA 2002)	$O(nm^{1/2})$	$O(n^3 T_c)=O(n^5)$	$O(m^{1/2})$
	3-HOP (Yang Xiang et al., SIGMOD '09)	$O(\log n + k)$	$O(kn^2)$	$O(nk)$

Depth First Traversal

- ? $u \rightsquigarrow v$
- Depth First Traversal (DFT) starting from u
- if node v is discovered:
 - then stop search, report YES
- If all nodes have been visited:
 - then report NO


No index and thus no
construction overhead
and no extra space
consumption

TOO GOOD

Query time: $O(m+n)$ the
entire graph should be
traversed in the worst
case

TOO BAD

Transitive Closure (TC)



	0	1	2	3	4	5
0	1	1	0	1	1	1
1	0	1	0	0	1	1
2	0	1	1	0	1	1
3	0	1	0	1	1	1
4	0	0	0	0	1	1
5	0	0	0	0	0	1

- It can be done by dynamic programming algorithm *Floyd–Warshall* in $\Theta(n^3)$
- It takes $O(n^2)$ space


TOO BAD

- BUT, queries can be answered in constant time $O(1)$

TOO GOOD

Optimal Tree Cover [idea]

SIGMOD '89



- Index: **smallest** postorder number of descendants

- DFT **postorder** number


Postorder: left-right-root

$label(u) = [u_{start}, u_{end}]$

(example) $?(b \rightsquigarrow h) \Rightarrow ?(1 \leq 2 < 4) \Rightarrow \text{YES}$
 $?(b \rightsquigarrow c) \Rightarrow ?(1 \leq 5 < 4) \Rightarrow \text{NO}$

Query Processing: $?(u \rightsquigarrow v) \Rightarrow$
 $?(u_{start} \leq v_{end} < u_{end})$

Optimal Tree Cover



Optimal Tree Cover

Topological
sorting: parent
comes before the
child

- **Index Construction:** $O(nm)$ time = $O(n^3)$
 1. Connect all nodes with no predecessors to a dummy node-root.
 2. Find a spanning tree for the DAG (+root) – requires a **topological sorting** first
 3. Label nodes according to tree edges (postorder).
 4. For each non-tree edge (u,v) – in **reverse topological order** of the nodes- :
$$label(v) = label(v) \cup label(u)$$
 5. If $label(a) \subset label(b)$ for labels inherited from nodes a, b , then keep only $label(b)$.


Optimal Tree Cover

- **Query processing:** $O(n)$ time

$$\text{label}(u) = \{[u_{\text{start}}, u_{\text{end}}], [u_{\text{start}_1}, u_{\text{end}_1}] \dots\}$$

$$(u \rightsquigarrow v) \text{ iff } \exists i: (u_{\text{start}_i} \leq v_{\text{end}} < u_{\text{end}_i})$$

- **Space:** $O(n^2)$ worst case



Bipartite Graph:
every one of these $(n-x)$
vertices will inherit all x labels,
resulting in total $(n-x)(x+1)$
labels, yielding $O(n^2)$.


Optimal Tree Cover

- The authors propose an algorithm for finding the **optimal spanning tree** (in terms of total label size).
- They also suggest a **maintenance mechanism** with **non-consecutive numbering** of nodes.
- Although **asymptotically equivalent** to the straightforward method of **transitive closure**, in the experiments the method performs better in **orders of magnitude**.
- The experiments are quite primitive (graphs with 1000 nodes)

GRIPP Index Creation


SIGMOD '07

- Depth-first traversal of G



- We reach a **node v**
 - for the first time
 - add *tree instance* of v to $IND(G)$
 - proceed traversal
 - **again**
 - add *non-tree instance* of v to $IND(G)$
 - do not traverse child nodes of v

GRIPP Index Table, $IND(G)$



Graph, G

node	pre	post	inst
R	0	21	tree
A	1	20	tree
B	2	7	tree
E	3	4	tree
F	5	6	tree
C	8	9	tree
D	10	19	tree
G	11	14	tree
B'	12	13	Non-tree
H	15	18	tree
A'	16	17	Non-tree

GRIPP index, $IND(G)$

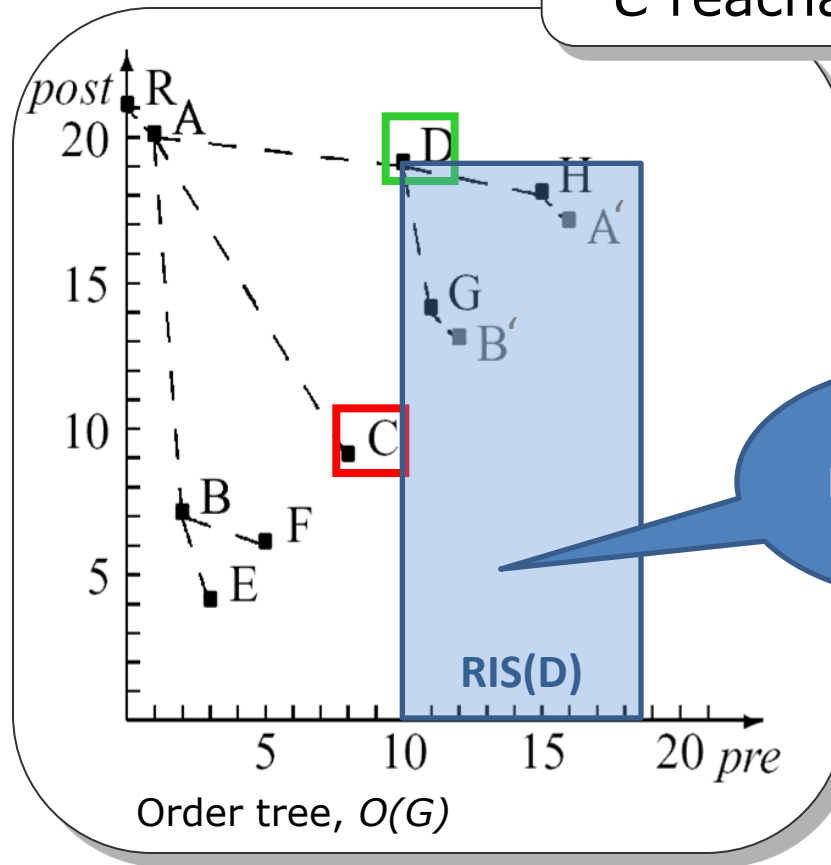
- Is node C reachable from node D?

GRIPP Query answering

$?(D \rightsquigarrow C)$

node	pre	post	inst
R	0	21	tree
A	1	20	tree
B	2	7	tree
E	3	4	tree
F	5	6	tree
C	8	9	tree
D	10	19	tree
G	11	14	tree
B'	12	13	non
H	15	18	tree
A'	16	17	non

RIS(D)



If $D_{pre} < C_{pre} < D_{post}$

C reachable from D

Reachable
Instance Set
of D

GRIPP Query answering

$?(D \rightsquigarrow C)$


Step 1: Retrieve RIS (D).

If $C \in RIS(D)$ then answer **YES** and finish.

Step 2: Else

foreach non_tree entry $h' \in RIS(D)$ do:

recursively issue the query $?(h' \rightsquigarrow C)$



h is the
correspondent
tree entry of h'

Gripp [Facts]


- **Index Construction Time:** Depth First Traversal, linear $O(m+n)$
- **Storage:** $O(n)$ nodes of the graph + $O(m-n)$ non-tree nodes yields $O(m+n)$ storage
- **Query Time:** in the worst case we ask for $O(m-n)$ recursive calls.

The authors use some **pruning techniques** and **heuristics** and claim that their algorithm has **almost constant** query time for various types of graphs.

The order of the traversal is crucial. The same for the order of the recursive hops.

Dual Labeling [assumptions]

- Basic assumption: most practical graphs are **sparse**.
 - Examples of biological data and XML documents
 - Average degree ~ 1.2 (*edges/node*)
- The authors will use this fact to build nearly optimal algorithms for **tree-like** sparse graphs
- Thus, they assume that $t \ll n$.




number of
non-tree
edges

Dual Labeling [idea]

ICDE '06

1. DFT to Compute a spanning tree and the Transitive Link Table (for the non-tree edges)
2. Compute the **transitive link closure**.



TRANSITIVE LINK TABLE


7 → [2,5)

9 → [6,9)

9 → [2,5)

Dual Labeling

- Now we can answer queries by checking the **tree labels** + the **transitive link table(TLT)**
 - examples: $?A \rightsquigarrow G$ **YES**: $7 \in [1,9)$
 - $?D \rightsquigarrow F$ **YES**: Although $4 \notin [6,9)$ if we **search** TLT we find edge $7 \rightarrow [2,5)$ for which $7 \in [6,9)$ and $4 \in [2,5)$
 - $?D \rightsquigarrow C$ **NO**: $5 \notin [6,9)$ and there is no entry in TLT with the above property



TRANSITIVE LINK TABLE

7→[2,5)

9→[6,9)

9→[2,5)

Dual Labeling

- The size of TLT is $O(t^2)$ since it contains the transitive closure of t non-tree edges.
- Given the above indexing scheme, query time might take $O(t^2)$ for the linear search of TLT
- The goal is to reduce query time to $O(1)$.
- We wouldn't mind to put them in a table (since t is small) in order to reduce the query time in $O(1)$ but we cannot!
 - TLT consists of entries of the form: $i \rightarrow [x, y)$
 - We would need 3D table
- The authors propose one solution

Dual Labeling

$?u \sim v$


$u = [a1, b1)$

$v = [a2, b2)$

$a2 \notin [a1, b1)$ so unreachable from tree only edges

What is the property of an entry $i \rightarrow [j, k)$ in TLT?

$i \in [a1, b1) \wedge [j, k) \ni a2$




Dual Labeling

TRANSITIVE LINK TABLE

7->[2,5)

9->[6,9)


9->[2,5)



TLT Grid

1	1	1
0	2	1
$\frac{y}{x}$	0	1

TLT Matrix



- $x = index_x(a')$, where $a' = \min\{i \mid i \rightarrow [j, k) \in T \wedge i \geq a\}$. If such an a' does not exist, let x be the special symbol “-.”
- $y = index_x(b')$, where $b' = \min\{i \mid i \rightarrow [j, k) \in T \wedge i \geq b\}$. If such a b' does not exist, let y be “-.”
- $z = index_y(a^*)$, where a^* is the start interval label of the lowest (tree) ancestor of u with a non-tree incoming edge. If such an a^* does not exist, let z be “-.”

Dual Labeling

Now reachability can be defined in constant time by the values of the two labels.

$u: ([a1, b1), < x1, y1, z1 >)$

$v: ([a2, b2), < x2, y2, z2 >)$

$u \rightsquigarrow v \Leftrightarrow$

- $a2 \in [a1, b1)$ or
- $N[x1, z2] - N[y1 - z2] > 0$

e.g.:

? $K \rightsquigarrow E$

- $3 \notin [9, 10)$
- $N[1, 0] - N[-, 0] = 1 - 0 > 0$

So the answer is **YES**.


? $H \rightsquigarrow B$

- $2 \notin [8, 9)$
- $N[1, 0] - N[1, 0] = 1 - 1 \ngtr 0$

So the answer is **NO**.

1	1	1
0	2	1
$\frac{y}{x}$	0	1

TLT Matrix



Dual Labeling [sum up]


- **Index Construction Time:** Depth First Traversal, linear $O(m+n)$ + transitive link closure construction $O(t^3)$ yields $O(m+n+t^3) \approx O(m+n)$ for $t \ll n$.
- **Storage:** $O(n)$ nodes of the graph $O(t^2)$ for the TLT matrix yields $O(n + t^2)$ storage
- **Query Time:** $O(1)$

Of course if t is comparable to n (there are a lot of non-tree edges) then Dual Labeling performs as bad as the naïve approach of the Transitive closure of the graph.

Chain Cover

ACM Trans. Database Syst. '90

- Enough with the spanning trees!
- Let's partition the graph into **chains**



Chain 0: $R \rightsquigarrow K \rightsquigarrow \cancel{D} \rightsquigarrow G \rightsquigarrow \cancel{B} \rightsquigarrow E$

Chain 1: $D \rightsquigarrow H$

Chain 2: $A \rightsquigarrow C$

Chain 3: $B \rightsquigarrow F$


$R \rightsquigarrow K \rightsquigarrow G \rightsquigarrow E$

Chain Cover


- A chain cover of G is a partition of G into **disjoint sets** called chains.
- Let $G=(V,E)$ a directed graph and $c_i \subseteq V$ s.t. if $u, v \in c_i$ then $u \rightsquigarrow v$. Now let $C=\{c_1, c_2, \dots, c_k\}$ the set containing such sets.

If $\forall u \in V \exists i: u \in c_i$ and $\forall i \neq j, c_i \cap c_j = \emptyset$ then C is a chain cover of G .

Chain Cover [index]



The idea behind the chain cover is again to produce a **compressed transitive closure** of the graph based on the chains.



e.g

? $K \rightsquigarrow F$

- Find label of F: (1,3)
- Find K's entry for chain no 3: (0,3)
- $0 \leq 1$ so answer is

YES


	R	A	B	C	D	E	F	G	H	K
c_0	(1,0)	(2,0)	(3,0)	-	(2,0)	-	-	(3,0)	-	(2,0)
c_1	(0,1)	(0,1)	-	-	(1,1)	-	-	-	-	(0,1)
c_2	(0,2)	(2,1)	-	-	-	-	-	-	-	-
c_3	(0,3)	(0,3)	(1,3)	-	(0,3)	-	-	(0,3)	-	(0,3)

Chain Cover

- The efficiency depends heavily in the initial chain covering (not unique of course).
- The smaller the number of chains the better.
- Optimal chain cover can be found in polynomial time.
- How? Transform the problem to a *min-flow* problem.

Find optimal Chain Cover

- We'll use a simpler graph for illustration:




Min-Flow problem

- Solve the **min flow** problem for the transformed graph under the constraints:
 - flow $(x_i y_i) > 0 \forall i$
 - flow $(x_i y_j) \geq 0 \forall i \neq j$
 - no flow accumulation in the nodes


The problem can be formulated by LP (Linear Program)

It can be solved Ford Fulkerson's algorithm



Chain Cover [facts]

- **Index Construction Time:** It takes $O(n^3)$ to compute the transitive closure and find the min chain cover
 - there are faster (and a lot more complicated) methods that use bipartite matching and drop the complexity to $O(n^2 + kn\sqrt{k})$
- **Storage:** $O(nk)$ [worst case: $k=O(n)$ so there is no real compression]
- **Query Time:** $O(1)$ if k is small enough to store the index in a 2D table. If not storing it into lists and indexing the lists yields $O(\log n + k)$ query time



k is the number of the chains

Path-Tree Cover


SIGMOD '08

- In the tree covering approaches, we tried to:
 1. build a spanning tree T ,
 2. give some labels w.r.t. T and then
 3. find a solution for the extra reachability induced by the non-tree edges
- In this paper the authors pay special attention to the **first** of the above steps.
- They generalize the notion of spanning tree to **spanning graph**
- They try to compute the **best** spanning graph in order to reduce the complexity of the third step (the non-spanning edges)

Constructing Path-Tree

- Step 1: Path-Decomposition of DAG
- Step 2: Minimal Equivalent Edge Set between any two paths
- Step 3: Path-Graph Construction
- Step 4: Path-Tree Cover Extraction

Step 1: Path-Decomposition





(PID, SID)
 $= (2, 5)$

For any two nodes (u, v)
in the same path,
 $u \rightarrow v$ if and only if $(u.sid \leq v.sid)$

Simple linear algorithm based on topological sort can achieve a path-decomposition


Step 2: Minimal equivalent edge set

The reachability between any two paths can be captured by a unique **minimal** set of edges




The edges in the minimal equivalent edge set do not cross (always parallel)!

Step 3: Path-Graph Construction




Weight reflects the cost we have to pay for the transitive closure computation if we exclude this path-tree edge




Weighted Directed Path-Graph

Step 4: Extracting Path-Tree Cover





Weighted Directed Path-Graph



Maximal Directed Spanning Tree


Chu-Liu/Edmonds algorithm, $O(m' + k \log k)$

3-Tuple Labeling for Reachability



DFS labeling (1-tuple)

Transitive Closure Compression



Path-tree cover (including labeling)
can be constructed in $O(m + n \log n)$

An efficient procedure can compute and compress the transitive closure in $O(mk)$, k is number of paths in path-tree


Path-tree cover [and then?]

- After building this complex index structure, they use the techniques discussed above (GRIPP, Dual Labeling etc.).

2 HOP-Cover [idea]

SODA '02


- Transitive closure compression (again...)
- If I choose node 1 as a **center node**, then I know that all 1's ancestors can reach 1's descendants.
- So by choosing node 1, all the green edges are covered.
- The goal is to choose nodes so as to cover **all edges in TC(G)**



2 HOP-Cover [idea]

SODA '02

- Based on that, we can **label nodes** as follows:
 - each node u will have a label $L(u)$
 - $L(u) = \{L_{in}(u), L_{out}(u)\}$
 - $L_{in}(u), L_{out}(u) \subseteq V$
- After choosing node 1, we add it at
 - $L_{out}(2) = \{1\}, L_{out}(3) = \{1\}, L_{out}(0) = \{1\}$
 - $L_{in}(4) = \{1\}, L_{in}(5) = \{1\}$
 - $L_{in}(1) = \{1\}, L_{out}(1) = \{1\}$
- After covering all edges of $TC(G)$, nodes are labeled



	0	1	2	3	4	5
L_{in}		{1}		{0}	{1,4}	{1,4}
L_{out}	{1,4,0}	{1}	{1}	{1,4}	{4}	

2 HOP-Cover [idea]

SODA '02

- Now reachability queries can be answered using the labels:

– $?u \rightsquigarrow v$

$$L_{out}(u) \cap L_{in}(v) \neq \emptyset$$

e.g.

$?0 \rightsquigarrow 5$

$$L_{out}(0) \cap L_{in}(5) = \{1,4,0\} \cap \{1,4\} \neq \emptyset$$

YES

$?4 \rightsquigarrow 1$

$$L_{out}(4) \cap L_{in}(1) = \{4\} \cap \{1\} = \emptyset$$


NO

Table with redundancy

Compression for each center node w

$$TC_{size} = Ans_w \cdot Desc_w$$

$$2\text{-HOP} = Ans_w + Desc_w$$



	0	1	2	3	4	5
L_{in}		{1}		{0}	{1,4}	{1,4}
L_{out}	{1, 4 , 0}	{1}	{1}	{1,4}	{4}	

2 HOP-Cover

Compression for each center node w

$$TC_{size} = Ans_w \cdot Desc_w$$

$$2\text{-HOP} = Ans_w + Desc_w$$

- **Problem:** How do you find a **minimum** 2-Hop Cover in the graph?
 - exact solution is NP hard
- **Approximate Solution:** Greedily pick the node with the **highest compression** as a center node
 - $\log n$ approximation ratio
 - **but** for each node we should compute all the subsets of ancestors and descendants to see which yields the highest compression
 - **exponential combinations:** another 2-approximation algorithm to find the approximate highest compression node (algorithm based on bipartite matching- works in linear time)

2-HOP cover

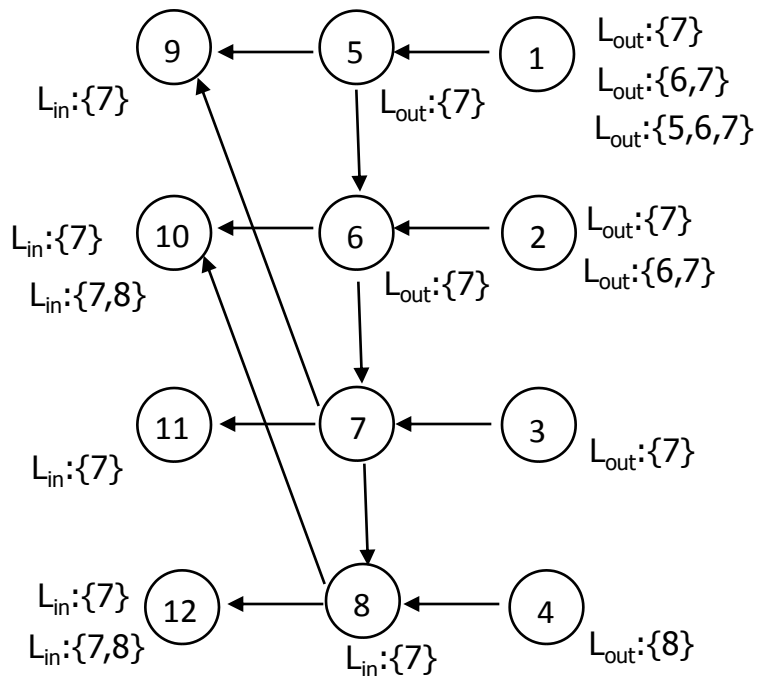
- **Problem:** How do you find a **minimum** 2-Hop Cover in the graph?
 - exact solution is NP hard
- There have been also other methods:
 - Heuristics
 - Geometrical Approach
 - Graph partitioning methods

2-HOP cover

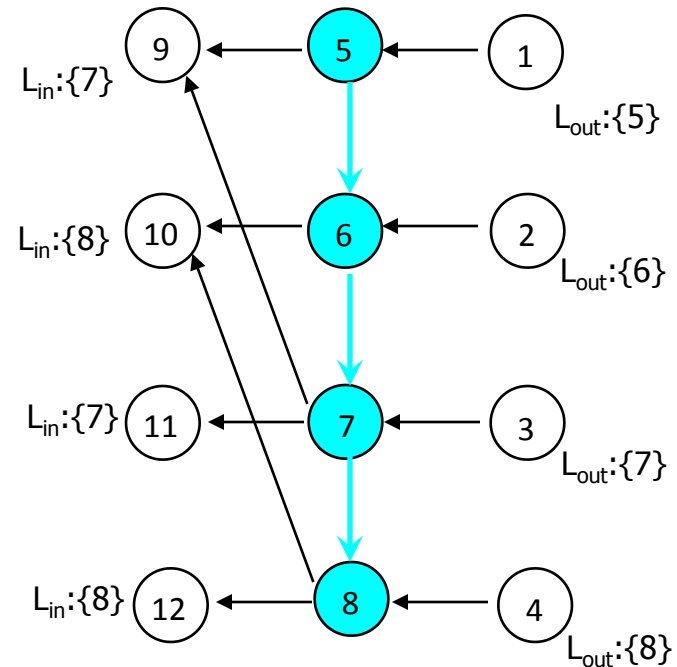
- **Index Construction Time:** complicated and costly $O(n^3 \cdot |TC|) = O(n^5)$
- **Storage:** $O(n\sqrt{m})$
- **Query Time:** $O(\sqrt{m})$.

3 HOP-Cover [intuition]

SIGMOD '09



2-Hop



3-Hop

3-HOP Cover [overview]

- Vertex \rightarrow Vertex \rightarrow Vertex (2 hop)
- Vertex \rightarrow Chain(Vertex \rightarrow Vertex) \rightarrow Vertex (Initial motivation of 3-HOP)
- Chain(Vertex \rightarrow Vertex) \rightarrow Chain(Vertex \rightarrow Vertex) \rightarrow Chain(Vertex \rightarrow Vertex) (3 hop contour)
- Chain decomposition is a spanning structure of G
- Some special vertices in the graph are labeled by L_{out} (a subset of vertices it can reach) and/or L_{in} (a subset of vertices it can be reached from).
- Chain decomposition plus the set of L_{out} and L_{in} are all that we need to design efficient reachability answering schemes.

Conclusion

	Method	Query time	Construction	Index size
Naïve	DFS/BFS	$O(n+m)$	$O(n+m)$	$O(n+m)$
	Transitive Closure	$O(1)$	$O(nm)=O(n^3)$	$O(n^2)$
Tree Cover	Optimal Tree Cover (Agrawal et al., SIGMOD'89)	$O(n)$	$O(nm)=O(n^3)$	$O(n^2)$
	GRIPP (Tripl et al., SIGMOD'07)	$O(m-n)$	$O(n+m)$	$O(n+m)$
	Dual-Labeling (Wang et al., ICDE'06)	$O(1)$	$O(n+m+t^3)$	$O(n+t^2)$
Chain Cover	Optimal Chain Cover (Jagadish, TODS'90)	$O(k)$	$O(nm)$	$O(nk)$
	Path-Tree (Jin, et al., SIGMOD'08)	$\log^2 k'$	$O(mk')/O(mn)$	$O(nk')$
HOP Cover	2-HOP (SODA 2002)	$O(nm^{1/2})$	$O(n^3 T_c)=O(n^5)$	$O(m^{1/2})$
	3-HOP (Yang Xiang et al., SIGMOD '09)	$O(\log n + k)$	$O(kn^2)$	$O(nk)$

THANK YOU

APPENDIX

Optimal Tree Cover[results]






Figure 3.9. Storage required for a 1000 node graph as a function of average degree

GRIPP[results]




(a) Average size (tuples)





(b) Average time (sec)

DFS labeling



1. Starting from the first vertex in the root-path
2. Always try to visit the next vertex in the same path
3. Label a node when all its neighbors has been visited
 $L(v) = N - x$, x is the # of nodes has been labeled

3-Tuple Labeling for Reachability



$u \rightarrow v$ if and only if 1) Interval label $I(u) \supseteq I(v)$
 2) DFS label $L(u) \leq L(v)$

?Query(9,15)

$P4[1,4] \supseteq P1[1,1]$ and $5 < 15$

Yes

?Query(9,2)

?Query(5,9)

3-HOP Cover [results]




Figure 8: Index size of Synthetic Datasets (2K)